

EECS 182 / 282A

Today: RNNs
Self-supervision

Architecture Order In Class:

MLPs \rightarrow CNNs \rightarrow Graph NN \rightarrow RNN/State-space \rightarrow Transformers

RNNs: Recurrent Neural Nets

History: Before 2018, consensus view: Two big successes in NN:
 Images & Vision using CNNs
 Language & Speech using RNNs

Key issue in language & speech (along with control, time-series, etc.) sequential d.h.

A signal-processing perspective

SP

Neural Nets

Weight-sharing

FIR Filter

(finite impulse response)

CNN

Across Space

IIR Filter

RNN

Across Time

(infinite impulse response)

Finite-dimensional hidden state inside filter.

Sequential & Causal


Inside an IIR filter from $\vec{u}_+ \rightarrow \vec{y}_+$, there is a hidden state \vec{h}_+ . (Initialized to zero)

e.g.

$$\begin{aligned}\vec{h}_{t+1} &= W_h \vec{h}_+ + B \vec{u}_{t+1}, \\ \vec{y}_{t+1} &= C \vec{h}_{t+1} + D \vec{u}_{t+1}\end{aligned}$$

(Update State Causally)


(Generate filter Output)



Treat this box as the counterpart of a conv for sequence problems.

Reading: Prince through Ch 11 + Ch 13
 Note: Little in this lecture is in the Prince textbook. (Just 9.3.7)
 Older RNN materials: deeplearningbook.org
Project Info Released: Start Forming Teams!

Approach: Keep all the design ideas & principles from CNNs.
Just replace space with time and enforce causality.



Many Ideas Just Carry Over
Vertical Residual Connections

Vertical " 1×1 " MLPs

Vertical Normalizations

Note: Can't average with the future
Just local channels

Key Design Question:

Where do the nonlinearities go?

Note: Use vertical direction for layers
Use horizontal for "time"

Traditional RNN perspective: Put a nonlinearity in the state update

Start with linear dynamics:

$$\begin{aligned}\vec{h}_{t+1} &= W_h \vec{h}_t + B \vec{u}_{t+1} + \vec{b}_h \\ \vec{y}_{t+1} &= C \vec{h}_{t+1} + \vec{b}_y\end{aligned}$$

biases

nonlinearity. typically sigmoid, tanh, etc.

Add a nonlinearity


$$\begin{aligned}\vec{h}_{t+1} &= \sigma(W_h \vec{h}_t + B \vec{u}_{t+1} + \vec{b}_h) \\ \vec{y}_{t+1} &= C \vec{h}_{t+1} + \vec{b}_y\end{aligned}$$

Linear Layer

Traditional "Vanilla" RNNs treat the hidden state as the layer's output.

A reason why traditional RNNs use saturating nonlinearities is that this blocks exploding activations across time.

Backprop during training:



So What's the Problem with Traditional RNNs?


Traditional Answer:

Had to learn long-range dependencies.

Many traditional answers and approaches... (LSTMs, GRUs, etc...)

Radical Approach: Think about implementation efficiency.

Parallelism during Training? 😐



Step Back for clarity: Think about Kalman Filters...

Underlying true linear process:

$$\vec{x}_{t+1} = A\vec{x}_t + B_u \vec{u}_{t+1} + \vec{w}_{t+1} \quad \text{where } \vec{w}_t, \vec{v}_t \text{ are iid } N(0, K_w)$$


$$\vec{y}_{t+1} = C\vec{x}_{t+1} + D\vec{u}_{t+1} + \vec{v}_{t+1} \quad \text{and } N(0, K_v)$$

If all we observe are \vec{u}_t , \vec{y}_t , and we know A, B, C, D, K_w, K_v , the optimal steady-state filter to estimate \vec{x}_t (in a mean-squared sense) has the form:

$$\begin{aligned}\hat{x}_{t+1} &= A\hat{x}_t + B_u \vec{u}_{t+1} + F(\vec{y}_{t+1} - [C(A\vec{x}_t + B_u \vec{u}_{t+1})] - D\vec{u}_{t+1}) \\ &= \tilde{A}\vec{x}_t + \tilde{B}\vec{u}_{t+1} + F(\vec{y}_{t+1})\end{aligned}$$


Predictable Parts

where $\tilde{A} = (I - FC)A$ and $F = P C^T (C P C^T + K_v)^{-1}$
 $\tilde{B} = (I - FC)B_u - FD$ when P solves D.A.R.E (Discrete Algebraic Riccati Equation)
 $P = A P A^T + K_w - A P C^T (C P C^T + K_v)^{-1} C P A^T$



But what if we don't know A, B, C, D, K_w, K_v ... And all we have is training data traces: $(\vec{x}_j^{[t]}, \vec{u}_j^{[t]}, \vec{y}_j^{[t]})_{t=1}^T$, where j is different such traces (possibly of different lengths)


- 1) Deem \tilde{A}, \tilde{B}, F learnable parameters. (Shared weights across time)
- 2) Put squared loss on $\hat{x}_j^{[t]}$ i.e. $\| \hat{x}_j^{[t]} - x_j^{[t]} \|_2^2$
- 3) Run with $u_j^{[t]}$, $y_j^{[t]}$ as inputs



Now, what if we only had traces $(\vec{u}_j^{(t)}, \vec{y}_j^{(t)})_{j=1}^J$ — No states. Just input & output

Question: Can we learn \tilde{A}, \tilde{B}, F for KF? Why or Why not?

No Can change coordinates for \vec{x} and this is unobservable from



This basically works! We learn the essence of the KF but not the specific coordinate system for \vec{x} . Practically, this means that we can learn the coordinate system with just a little bit of \vec{x} data.

General Principle: Self-supervision

"I need labels to train, I don't have labels, so make my own labels from data"

Lessons from example above:

- 1) We can learn a partial pattern that can be useful
- 2) Might need **scaffolding** parts of my NN.
- 3) Generic idea of "next-thing" prediction in causal sequence modeling

Step Back: Connect to unsupervised learning in classic ML.

Two approaches:

- 1) Dimensionality Reduction ← Start here
- 2) Clustering

Recall Dimensionality Reduction.

Think about PCA

All we have are $\{\vec{x}_i\}_{i=1}^N \leftarrow d\text{-dim}$ Unlabeled Data From Interesting Distribution

Classic Recipe (Neglecting Means):

1) Construct

$$X = \begin{bmatrix} \vec{x}_1^T \\ \vec{x}_2^T \\ \vdots \\ \vec{x}_N^T \end{bmatrix}$$

2) Compute SVD $X = U \sum V^T = \sum_i \sigma_i \vec{u}_i \vec{v}_i^T$

3) Keep top k $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k$ singular vectors to use for dim-reduct

4) Given some other problem $\vec{x} \rightarrow \begin{bmatrix} \vec{v}_1^T \vec{x} \\ \vdots \\ \vec{v}_k^T \vec{x} \end{bmatrix} \} k\text{-dim features}$

Why was this reasonable?

Recall Eckhart-Young-Mirsky Theorem for Frobenius Norm

Given X , $\hat{X} = \sum_{i=1}^k \sigma_i \vec{u}_i \vec{v}_i^T$ is the rank- k matrix that minimizes

$$\|X - \hat{X}\|_F^2$$