

Lecture 17

- State space models continued

→ How to support long-range dependancies?

$$\begin{aligned}\vec{h}_{t+1} &= A \vec{h}_t + B \vec{x}_t \\ \vec{y}_t &= C \cdot \vec{x}_t\end{aligned}\quad \left.\right\} \text{All information about the past can only be captured in } \vec{h}_t. \text{ It has to fit.}$$

Efficiency versus Efficacy tradeoff.

Larger $\vec{h}_t \Rightarrow$ more memory/history tracking

but also, more compute.

So want to make \vec{h}_t as large as possible, while still being able to run

S4

Efficiency via FFT - use of LTI systems.

What about long range dependency?

Let $A = \begin{bmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & 0 & & \lambda_N \end{bmatrix}$ It is N dimensional.

How is λ related to system stability?

$|\lambda| > 1 \rightarrow$ System explodes \rightarrow bad.

$|\lambda| < 1 \rightarrow$ history dies out fast.

$|\lambda| = 1 \rightarrow$ critical damping.

For $\lambda \in \mathbb{R}$, $|\lambda|=1 \Rightarrow \lambda = \pm 1$.

→ Only two choices!

→ Far apart from each other (So not GD friendly)

Solution: Complex eigenvalues!

$$\lambda = \exp(-\text{ReLU}(\lambda_R) + j\lambda_I) \quad j = \sqrt{-1}$$

- Exact initialization matters, have to be mostly on or very close to unit circle with some inside unit circle
- Hippo-based initialization - like analytic pretraining.
- Freezing A with good initialization and only training the rest also seemed to work reasonably!

Table 4: (**Long Range Arena**) (*Top*) Original Transformer variants in LRA. Full results in Appendix D.2. (*Bottom*) Other models reported in the literature. *Please read Appendix D.5 before citing this table.*

MODEL	LISTOPS	TEXT	RETRIEVAL	IMAGE	PATHFINDER	PATH-X	AVG
Transformer	36.37	64.27	57.46	42.44	71.40	✗	53.66
Reformer	<u>37.27</u>	56.10	53.40	38.07	68.50	✗	50.56
BigBird	36.05	64.02	59.29	40.83	74.87	✗	54.17
Linear Trans.	16.13	<u>65.90</u>	53.09	42.34	75.30	✗	50.46
Performer	18.01	65.40	53.82	42.77	77.05	✗	51.18
FNet	35.33	65.11	59.61	38.67	<u>77.80</u>	✗	54.42
Nyströmformer	37.15	65.52	<u>79.56</u>	41.58	70.94	✗	57.46
Luna-256	37.25	64.57	79.29	<u>47.38</u>	77.72	✗	<u>59.37</u>
S4	59.60	86.82	90.90	88.65	94.20	96.35	86.09

Outperformed Transformer models on long-range benchmarks!

But didn't quite match/beat transformers on other tasks.

Mamba / S6.

- Recurrent view
(state-based)

"inference time"

State-space models

- Input/Output view,
(convolution based)

"training time"

- Continuous-time view

"selectivity) gating"

Discrete-time system:

$$\vec{h}_{t+1} = A \vec{h}_t + B \vec{u}_t$$


$$\vec{y}_{t+1} = C \cdot \vec{h}_t$$

Think of as a discretization of a continuous time system.
(sampling)

$$\frac{d}{dt} \vec{h}(t) = A \vec{h}(t) + B \vec{u}(t) \quad \leftarrow (\text{Different } A, B)$$

$$\vec{y}(t) = C \cdot \vec{h}(t)$$

To convert from CT to DT, we need a sampling interval ΔT .




Scalar case:

$$\dot{x}(t) = ax(t) + b \cdot u(t). \quad \text{Sampling time } \Delta T.$$

Assume $u(t)$ is piecewise constant for time ΔT .

For small ΔT this is reasonable.



Aside:

$a > 0 \Rightarrow \text{explode}$

$a < 0 \Rightarrow \text{decay to 0.}$

$$x(\Delta T) = e^{a\Delta T} \cdot x(0) + \underbrace{\frac{b}{a} (e^{a\Delta T} - 1)}_{\text{impact of input}} u(0).$$

$e^{a\Delta T}$ exponential growth.

Assume $a < 0$

Then ΔT close to 0 $\Rightarrow e^{a\Delta T} \approx 1 \Rightarrow$ initial state? Remember history.

ΔT large $\Rightarrow e^{a\Delta T}$ decaying exponential \Rightarrow forget past

$$\Delta T \text{ smallish} \Rightarrow e^{a\Delta T} - 1 \approx \underbrace{1 + a\Delta T - 1}_{\text{Taylor expansion.}} \approx a\Delta T.$$

$$\Rightarrow \frac{b}{a} (e^{a\Delta T} - 1) = b\Delta T.$$

ΔT close to 0 $\Rightarrow b\Delta T \approx 0$.

Effect of input?
input doesn't matter

ΔT large $\Rightarrow b\Delta T$ large \Rightarrow

input matters a lot.

So by changing ΔT , can change the relative importance of the 'past'.

Idea: Choose ΔT based on input (i.e. learn it).


Try: Linear function of \vec{w} ?

$$\langle \vec{w}, \vec{u} \rangle$$

Try: ReLU ($\langle \vec{w}, \vec{u} \rangle + b$)

Softplus:

$$\text{Softplus}(x) = \ln(1+e^x)$$



if $a=b=N=1$

$$g_t = \sigma(\text{Linear}(\vec{x}_t))$$

$$\vec{h}_t = g_t \vec{h}_{t-1} + (1-g_t) \vec{x}_t$$

So this gives us for each channel of input \vec{x} :

$$\vec{h}_{k+1} = A_k \vec{h}_k + B_k \vec{x}_k$$

$$\vec{y}_{k+1} = C \vec{h}_k$$

$$A_k = \begin{bmatrix} e^{\lambda_1 \Delta T} & & \\ & e^{\lambda_2 \Delta T} & \\ & & \ddots \end{bmatrix}$$

But... have we lost anything?

$$B_k = \vec{b}_k = \vec{b} \Delta T.$$

$$A^k \rightarrow A_k A_{k-1} \cdots A_1 A_0$$

Initialization:

λI blocks for A.

$$\lambda_n = -\frac{1}{2} + ni$$

$$\lambda_n = -\frac{1}{2}$$

$$\lambda_n = -(n+1)$$

State-space models

- Recurrent view
(state-based)


"inference time"

- Input/Output view
(Convolution based)

"training time"

- Continuous-time View

"selectivity) gating"



Sequential Mamba Block